



GNU make

A Program for Directing Recompilation

Auto-Dependency Generation

June 23rd, 2000 (updated April 9th, 2017)

One of the most important, and yet potentially frustrating, tasks that is required to allow any `make`-based build environment to function properly is the correct listing of dependencies in the makefile.

This document describes a very useful method for having `make` itself create and maintain these dependencies completely automatically.

The basics of the advanced method below were invented by Tom Tromey <tromey@cygnus.com>, I merely wrote it down here. Credits for the method go to him; problems with the explanation belong to me.

- [TL;DR: The GCC Solution](#)
- [Traditional `make depend` Method](#)
- [The GNU `make include` Directive](#)
- [Basic Auto-Dependencies](#)
- [Advanced Auto-Dependencies](#)
 - [Avoiding Re-exec of `make`](#)
 - [Avoiding "No rule to make target ..." Errors](#)
 - [Handling Deleted Dependency Files](#)
- [Placement of Output Files](#)
- [Defining `MAKEDEPEND`](#)
 - `MAKEDEPEND = /usr/lib/cpp`
 - `MAKEDEPEND = makedepend`
 - `MAKEDEPEND = gcc -M`
- [Combining Compilation and Dependency Generation](#)
- [Dependencies For Non-C Files](#)

All make programs must know, with great accuracy, what files a particular target is dependent on in order to ensure that it is rebuilt when (and only when) necessary.

Keeping this list up-to-date by hand is not only tedious, but quite error prone. Most systems of any size prefer to provide automated tools for extracting this information. The traditional tool was the `makedepend` program, which reads C source files and generates a list of the header files in a target-dependency format suitable for inclusion in (or appending to) a makefile.

The modern solution, for those using suitably enlightened compilers or preprocessors (such as GCC), is to have the compiler or preprocessor generate this information.

The purpose of this paper isn't primarily to discuss ways in which this dependency information can be obtained, although this is covered. Rather, it describes some useful ways to combine the invocation and output of these tools with GNU `make` to ensure that dependency information is always accurate and up-to-date, as seamlessly (and efficiently) as possible.

These methods rely on features provided by GNU `make`. It may be possible to modify them to work with other versions of `make`; that's left as an exercise to the reader. However, before undertaking that exercise please review [Paul's First Rule of Makefiles](#) :).

TL;DR: The GCC Solution

For those who are impatient, here's a complete best-practice solution. This solution requires support from your compiler: it assumes you are using [GCC](#) as your compiler (or a compiler which provides preprocessor flags compatible with GCC). If your compiler doesn't meet this criteria, keep reading for alternatives.

Add this to your makefile environment (the sections in blue are the changes to the built-in content provided by GNU `make`). Of course you can omit pattern rules which don't fit your environment (or add new ones as needed):

```
DEPDIR := .d
$(shell mkdir -p $(DEPDIR) >/dev/null)
DEPFLAGS = -MT $@ -MMD -MP -MF $(DEPDIR)/$*.Td

COMPILE.c = $(CC) $(DEPFLAGS) $(CFLAGS) $(CPPFLAGS) $(TARGET_ARCH) -c
COMPILE.cc = $(CXX) $(DEPFLAGS) $(CXXFLAGS) $(CPPFLAGS) $(TARGET_ARCH) -c
POSTCOMPILE = @mv -f $(DEPDIR)/$*.Td $(DEPDIR)/$*.d && touch $@

%.o : %.c
%.o : %.c $(DEPDIR)/%.d
    $(COMPILE.c) $(OUTPUT_OPTION) $<
    $(POSTCOMPILE)

%.o : %.cc
%.o : %.cc $(DEPDIR)/%.d
    $(COMPILE.cc) $(OUTPUT_OPTION) $<
    $(POSTCOMPILE)
```

```

%.o : %.cxx
%.o : %.cxx $(DEPDIR)/%.d
    $(COMPILE.cc) $(OUTPUT_OPTION) $<
    $(POSTCOMPILE)

$(DEPDIR)/%.d: ;
.PRECIOUS: $(DEPDIR)/%.d

include $(wildcard $(patsubst %, $(DEPDIR)/%.d, $(basename $(SRCS))))

```

Be aware that the `include` line must come *after* the initial, default target; otherwise the included dependency files will abscond with your default target setting. It's fine to add this to the very end of your makefile (or, keep it in a separate makefile and include it).

Also, this assumes the `SRCS` variable contains all the source files (not header files) you want to track dependencies for.

If you just want to understand what these changes mean, [jump to the description](#). If you want a deeper understanding of how this works, and to learn some interesting things about GNU `make`, keep reading.

Traditional `make depend` Method

The time-honored method of handling dependency generation is to provide a special target in your makefiles, typically `depend`, which can be invoked to create dependency information. The command for this target will invoke some kind of dependency tracking tool on all the relevant files in the directory, which will generate makefile-formatted dependency information.

If your version of `make` supports `include` you can redirect this to a file and simply include the file. If not, this usually also involves some shell hackery to append the list of dependencies generated to the end of the makefile itself.

Although it's simple, there are serious problems with this method. First and foremost is that dependencies are only rebuilt when the user explicitly requests it; if the user doesn't run `make depend` regularly they could become badly out-of-date and `make` will not properly rebuild targets. Thus, we cannot say this is seamless and accurate.

A secondary problem is that it is inefficient to run `make depend`, particularly after the first time. Since it modifies makefiles, you typically must do it as a separate build step, which means an extra invocation of every `make` in every subdirectory, etc., in addition to the overhead of the dependency-generation tool itself. Also, it rechecks every file's dependencies, even those which haven't changed.

So, we'll see how we can do better.

The GNU `make include` Directive

Most versions of `make` support some sort of `include` directive (and indeed, `include` is required by the latest POSIX specification). Unsurprisingly, this allows one makefile to include other makefiles, as if they had been entered there.

One can immediately see how this would be useful, simply to avoid appending dependency information to a makefile as in the step above. However, there is a more [interesting capability](#) in GNU `make`'s handling of `include`: just as with the normal makefile, GNU `make` will attempt to *rebuild* the included makefile. If it is successfully rebuilt, GNU `make` will re-execute itself to read the new version.

This auto-rebuild feature can be harnessed to avoid requiring a separate `make depend` step: if you list all the source files as prerequisites to the file containing dependency information, then include that file into your makefile, it will be rebuilt every time a source file changed. As a result, the dependency information will always be up-to-date and the user doesn't need to run `make depend` explicitly.

Of course, this means dependency information is recalculated for *all* files every time *any* file changes, which is unfortunate. We can still do better than this.

For a detailed description of GNU `make`'s automatic rebuild feature, see the GNU `make` User's Manual, section "*How Makefiles Are Remade*".

Basic Auto-Dependencies

The GNU `make` User's Manual describes one way of handling auto-dependencies in section [Generating Dependencies Automatically](#).

In this method, a separate dependency file is created for each source file (in our examples we'll use a `.d` suffix on the base filename). This file contains a rule for the target that is created from that source file, listing the generated prerequisites of the target.

These dependency files are then all included by the makefile. An implicit rule is provided that describes how the dependency files are to be created. In short, something like this:

```
SRCS = foo.c bar.c ...

%.d : %.c
    $(MAKEDEPEND)
    @sed 's/\($*\)\.o[ :]*\/1.o $@ : /g' < $*.Td > $@; \
        rm -f $*.Td; [ -s $@ ] || rm -f $@

include $(SRCS:.c=.d)
```

In these examples I'll simply use variables such as `$(MAKEDEPEND)` to stand for whatever method you choose for creating dependency files. Some possible values for this variable are described below.

In this case, the output is written to a temporary file and post-processed. The post-processing changes the normal target specification:

```
foo.o: foo.c foo.h bar.h baz.h
```

to also contain the `.d` file itself, like this:

```
foo.o foo.d: foo.c foo.h bar.h baz.h
```

Whenever GNU `make` reads this makefile, before it does anything else it will try to rebuild the included makefiles, in this case the `.d` files. We have a rule to build them, and we have a list of prerequisites which is the same as the list for the `.o` file itself. So, if any file changes that would cause the original target to appear out-of-date, it will also cause the `.d` file to be rebuilt.

Thus, when any source file or included file changes, `make` will rebuild the `.d` file(s), re-execute itself to read in the new makefiles, then continue with builds as usual, now with an updated and accurate dependency list.

Here we solve the two problems with the earlier solutions. First, the user doesn't have to do anything special to ensure accurate dependency lists; `make` will take care of it. Second, we are only updating dependency lists for those file which have actually changed, not all the files in the directory.

We have three new problems with this method, however. The first is still efficiency. Although we only re-examine changed files we still will re-exec `make` if anything changes, which could be slow for large build systems.

The second problem is merely annoying: when you add a new file or build for the first time, no `.d` file will exist. When `make` tries to include it and discovers it doesn't exist, it will generate a warning. GNU `make` will then proceed to rebuild the `.d` file and re-invoke itself, so this is not fatal; nevertheless it's annoying.

The third problem is more serious: if you remove or rename a prerequisite file (say a C `.h` file), `make` will stop with a fatal error, complaining that the target doesn't exist:

```
make: *** No rule to make target 'bar.h', needed by 'foo.d'. Stop.
```

This is because the `.d` file has a dependency on a file `make` can't find. It can't rebuild the `.d` file until all the prerequisites are there, and it won't realize it doesn't need the prerequisite until it rebuilds the `.d` file. Catch-22.

The only solution here is to go in by hand and remove any `.d` file that refers to the missing file—typically it's simpler to remove all of them than try to find the proper ones. You can even create a `clean-deps` target or similar to do it automatically (investigate the `MAKECMDGOALS` variable to see how you can write this target while still avoiding the rebuild attempt on the `.d` files). This is annoying to be sure, but given that files aren't removed or renamed all that often in a typical environment perhaps it's not fatal.

Advanced Auto-Dependencies

The basis for the method described here was engineered by Tom Tromey <tromey@cygnus.com>, who used it as the standard dependency generation method for the FSF's [automake](#) tool. I've made some changes to allow it to fit into a more generic build environment.

Avoiding Re-exec of `make`

Let's address the first problem above: the re-invocation of `make`. If you think about it, this re-invocation is really not necessary. Since we know *some* prerequisite of the target changed, we must rebuild the target; having a more up-to-date list won't affect that decision. What we really need is to ensure that the prerequisite list is up-to-date for the *next* invocation of `make`, when we need to decide whether to update it again.

Since we don't need the up-to-date prerequisite list in this build, we can actually avoid re-invoking `make` at all: we can simply have the prerequisite list built *at the same time* as the target is rebuilt. In other words, we can change the build rule for our target to add in commands to update the dependency file. Also, in this case we must be very careful that we *don't* provide rules to build the dependencies automatically: if we do `make` will still try to rebuild them and re-exec: we don't want that.

Now that we don't care about dependency files that don't exist, solving the second problem (superfluous warnings) is easy: we can just use GNU `make`'s `wildcard` function so that dependency files which don't exist won't cause an error.

Let's take a look at an example so far:

```
SRCS = foo.c bar.c ...

%.o : %.c
    @$(MAKEDEPEND)
    $(COMPILE.c) -o $@ $<

include $(wildcard $(SRCS:.c=.P))
```

Avoiding "No rule to make target ..." Errors

This one is a little trickier. However, it turns out we can convince `make` to not fail merely by mentioning the file explicitly as a target in the makefile. If a target exists, but has no commands (either implicit or explicit) or prerequisites, then `make` simply always considers it up-to-date. That's the normal case, and it behaves as we'd expect.

In the case where the above error occurs, the target *doesn't* exist. According to the GNU `make` User's Manual section [Rules without Recipes or Prerequisites](#):

If a rule has no prerequisites or recipe, and the target of the rule is a nonexistent file, then `make' imagines this target to have been updated whenever its rule is run. This implies that all targets depending

on this one will always have their recipes run.

Perfect. It ensures `make` won't throw an error since it knows how to handle that non-existent file, and it ensures that any file depending on that target is rebuilt, which is exactly what we want.

So, all we need to do is post-process the original dependency file and turn all the prerequisites into targets with no commands or prerequisites. Something like this [1]:

```
SRCS = foo.c bar.c ...

%.o : %.c
    @$(MAKEDEPEND); \
    cp $*.Td $*.d; \
    sed -e 's/#.*//' -e 's/^[^:]*: *//' -e 's/ *\\$$//' \
        -e '/^$$/ d' -e 's/$$/ :/' < $*.Td >> $*.d; \
    rm -f $*.Td
    $(COMPILE.c) -o $@ $<

include $(wildcard $(SRCS:.c=.d))
```

Briefly, this creates a `.d` file with the original prerequisites list, then adds targets to it by taking each line, removing any existing target information and any line continuation (`\`) characters, then adding a target separator (`:`) to the end. This works with the values for `MAKEDEPEND` I suggest below; it's possible you will need to modify the translation for other dependency generators you might use.

Handling Deleted Dependency Files

There is one remaining issue with this configuration: if a user happens to delete dependency files without modifying any of the source files, `make` will not notice anything amiss and will not recreate the dependency files until it decides to rebuild the corresponding object file for some other reason. In the meantime `make` will be missing dependency information for those targets (so, for example, modifying a header file without changing the source file will not cause the object file to be rebuilt).

The problem is slightly complex because we do *not* want the dependency files to be considered “real” targets: if they are then when we use `include` to include them, `make` will rebuild them then re-exec itself. This isn't the end of the world, but it's overhead we'd prefer to do without.

The automake method doesn't address this problem. In the past I've suggested a “just don't do that” solution, combined with placing dependency files [in a separate directory](#) to make it difficult to accidentally delete them.

However, [Lukas Waymann](#) suggested a tidy solution: add the dependency file as a *prerequisite* to the target, and create an empty recipe for it:

```
SRCS = foo.c bar.c ...
```

```

%.o : %.c %.d
    @$(MAKEDEPEND); \
    cp $*.Td $*.d; \
    sed -e 's/#.*//' -e 's/^[^:]*: *//' -e 's/ *\\$$/ ' \
        -e '/^$$/ d' -e 's/$$/ :/' < $*.Td >> $*.d; \
    rm -f $*.Td
    $(COMPILE.c) -o $@ $<

%.d: ;

include $(wildcard $(SRCS:.c=.d))

```

This solves the problem nicely: when `make` checks the target it will see the dependency file as a prerequisite and try to build that. If it exists nothing will be done since there's no prerequisites for the dependency file. If it doesn't exist, it will be marked out of date since it has an empty recipe, which will force the object target to be rebuilt (creating a new dependency file).

When `make` is trying to rebuild included files it will find the implicit rule for the dependency and use it. However since the rule doesn't update the target file, no included file was updated and `make` will not re-execute itself.

Placement of Output Files

You may decide you don't like all those `.d` files cluttering up your source directory. You can easily have your makefile put them somewhere else. Here's an example of doing that for the advanced method; you can extrapolate to the other methods:

```

DEPDIR = .deps
df = $(DEPDIR)/$(*)

SRCS = foo.c bar.c ...

%.o : %.c
    @$(MAKEDEPEND); \
    cp $(df).Td $(df).d; \
    sed -e 's/#.*//' -e 's/^[^:]*: *//' -e 's/ *\\$$/ ' \
        -e '/^$$/ d' -e 's/$$/ :/' < $(df).Td >> $(df).d; \
    rm -f $(df).Td
    $(COMPILE.c) -o $@ $<

include $(wildcard $(SRCS:%.c=$(DEPDIR)/%.d))

```

Replace any references to `$*.Td` in the `MAKEDEPEND` script with `$(df).Td`.

Defining MAKEDEPEND

Here I'll discuss some possible ways to define the `MAKEDEPEND` variable I've blithely been using above. These definitions create output files named `.Td` (temporary dependency), since we typically need to perform some post-processing on them before they are put into use.

```
MAKEDEPEND = /usr/lib/cpp
```

The most basic way to generate dependencies is by using the C preprocessor itself. This requires a bit of knowledge about the output format of your preprocessor—luckily most UNIX preprocessors have similar output for our purposes. In order to preserve line number information for the compiler's error messages and debugging information, the output of the preprocessor must provide line number and *filename* information for each jump to an `#include` file and each return from one. These lines can be used to figure out what files were included.

Most UNIX preprocessors insert special lines in the output with this format:

```
# lineno "filename" extra
```

All we care about is the *filename* value. If your preprocessor generates output like the above, a definition like this for `MAKEDEPEND` should work:

```
MAKEDEPEND = $(CPP) $(CPPFLAGS) $< \
    | sed -n 's/^#\s*[0-9][0-9]*\s*"([^"]*)"\s*.*$/\1/p' \
    | sort | uniq > $*.Td
```

If you're using the advanced method, you can replace the `$.o` in the `sed` script with `$@`. If you have a modern version of `sort`, you can also replace `sort | uniq` with just `sort -u`.

And, of course, if you go this route you might as well combine the post-processing involved with the method you're using right into this script.

```
MAKEDEPEND = makedepend
```

The X Windowing System source tree comes with a program called `makedepend`. This program examines C source and header files and generates make dependency lines. It's geared towards adding those dependencies to the bottom of an existing makefile, so to use it the way we want we need to be a little tricky. For example, some versions fail if the output file doesn't already exist.

This definition should suffice:

```
MAKEDEPEND = touch $*.Td && makedepend $(CPPFLAGS) -f $*.Td $<
```

```
MAKEDEPEND = gcc -M
```

The [GNU Compiler Collection](#) contains a C preprocessor that can generate make dependency files. This definition should suffice:

```
MAKEDPEND = gcc -M $(CPPFLAGS) -o $*.Td $<
```

However, if you're using GCC you might as well continue on to the next section for a simpler and faster alternative.

Combining Compilation and Dependency Generation

If you're using GCC (or a compiler which provides equivalent options) you can save yourself a lot of time during the build by *combining* the dependency generation and the object file generation into a single command, because the compiler has the ability to output the dependency information as a side-effect of the compilation. Here is an example implementation, copied from the *TL;DR* section at the top of this page:

```
DEPDIR := .d
$(shell mkdir -p $(DEPDIR) >/dev/null)
DEPFLAGS = -MT $@ -MMD -MP -MF $(DEPDIR)/$*.Td

COMPILE.c = $(CC) $(DEPFLAGS) $(CFLAGS) $(CPPFLAGS) $(TARGET_ARCH) -c
COMPILE.cc = $(CXX) $(DEPFLAGS) $(CXXFLAGS) $(CPPFLAGS) $(TARGET_ARCH) -c
POSTCOMPILE = @mv -f $(DEPDIR)/$*.Td $(DEPDIR)/$*.d && touch $@

%.o : %.c
%.o : %.c $(DEPDIR)/%.d
    $(COMPILE.c) $(OUTPUT_OPTION) $<
    $(POSTCOMPILE)

%.o : %.cc
%.o : %.cc $(DEPDIR)/%.d
    $(COMPILE.cc) $(OUTPUT_OPTION) $<
    $(POSTCOMPILE)

%.o : %.cxx
%.o : %.cxx $(DEPDIR)/%.d
    $(COMPILE.cc) $(OUTPUT_OPTION) $<
    $(POSTCOMPILE)

$(DEPDIR)/%.d: ;
.PRECIOUS: $(DEPDIR)/%.d

include $(wildcard $(patsubst %, $(DEPDIR)/%.d, $(basename $(SRCS))))
```

Let's go over this:

DEPDIR = ...

This implementation places dependency files into a subdirectory named `.d`.

```
$(shell mkdir -p $(DEPDIR) >/dev/null)
```

Unfortunately GCC will not create subdirectories, so this line ensures that the `DEPDIR` directory always exists.

```
DEPFLAGS = ...
```

These are the special GCC-specific flags which convince the compiler to generate the dependency file. Full descriptions can be found in [the GCC manual](#) section [Options Controlling the Preprocessor](#):

```
-MT $@
```

Set the name of the target in the generated dependency file.

```
-MMD
```

Generate dependency information as a side-effect of compilation, not instead of compilation. This version omits system headers from the generated dependencies: if you prefer to preserve system headers as prerequisites, use `-MD`.

```
-MP
```

Adds a target for each prerequisite in the list, to avoid errors when deleting files.

```
-MF $(DEPDIR)/$.Td
```

Write the generated dependency file to a temporary location `$(DEPDIR)/$.Td`.

```
POSTCOMPILE = ...
```

First rename the generated temporary dependency file to the real dependency file. We do this in a separate step so that failures during the compilation won't leave a corrupted dependency file.

Second touch the object file; it's been reported that some versions of GCC may leave the object file *older* than the dependency file, which causes unnecessary rebuilds.

```
%.o : %.c
```

Delete the built-in rules for building object files from `.c` files, so that our rule is used instead. Do the same for the other built-in rules.

```
... $(DEPDIR)/%.d
```

Declare the generated dependency file as a prerequisite of the target, so that if it's missing the target will be rebuilt.

```
$(DEPDIR)/%.d: ;
```

Create a pattern rule with an empty recipe, so that `make` won't fail if the dependency file doesn't exist.

```
.PRECIOUS: $(DEPDIR)/%.d
```

Mark the dependency files precious to `make`, so they won't be automatically deleted as intermediate files.

include ...

Include the dependency files that exist: translate each file listed in `SRCS` into its dependency file. Use `wildcard` to avoid failing on non-existent files.

Dependencies For Non-C Files

In general you need some way of producing dependency files in order to use these methods. If you're working with files that aren't C files you'll need to discover or write your own method. Anything that generates make dependency files will do. This usually isn't too difficult.

An interesting idea has been proposed by Han-Wen Nienhuys <hanwen@cs.uu.nl> and he has a small "proof of concept" [implementation](#), although it currently only works on Linux. He suggests using the `LD_PRELOAD` environment to insert special shared library that contains a replacement for the `open(2)` system call. This version of `open()` would actually write out make dependency information for every file the commands read during operation. This would give you completely reliable dependency information for every kind of command without needing any special dependency extraction tools at all. In his proof-of-concept implementation you can control the output file and exclude some kinds of files (shared libraries, maybe) via environment variables.

[1] Note I have modified the post-processing sed scripts here from what Tom uses in automake, in order to allow various styles of `MAKEDEPEND` output to work.
